

LINEAR REGRESSION MODEL ---> USA HOUSING PRICES

```
In [2]: #imports
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

In [4]: #Reading file
df = pd.read_csv('USA_Housing.csv')
df.head(2)
```

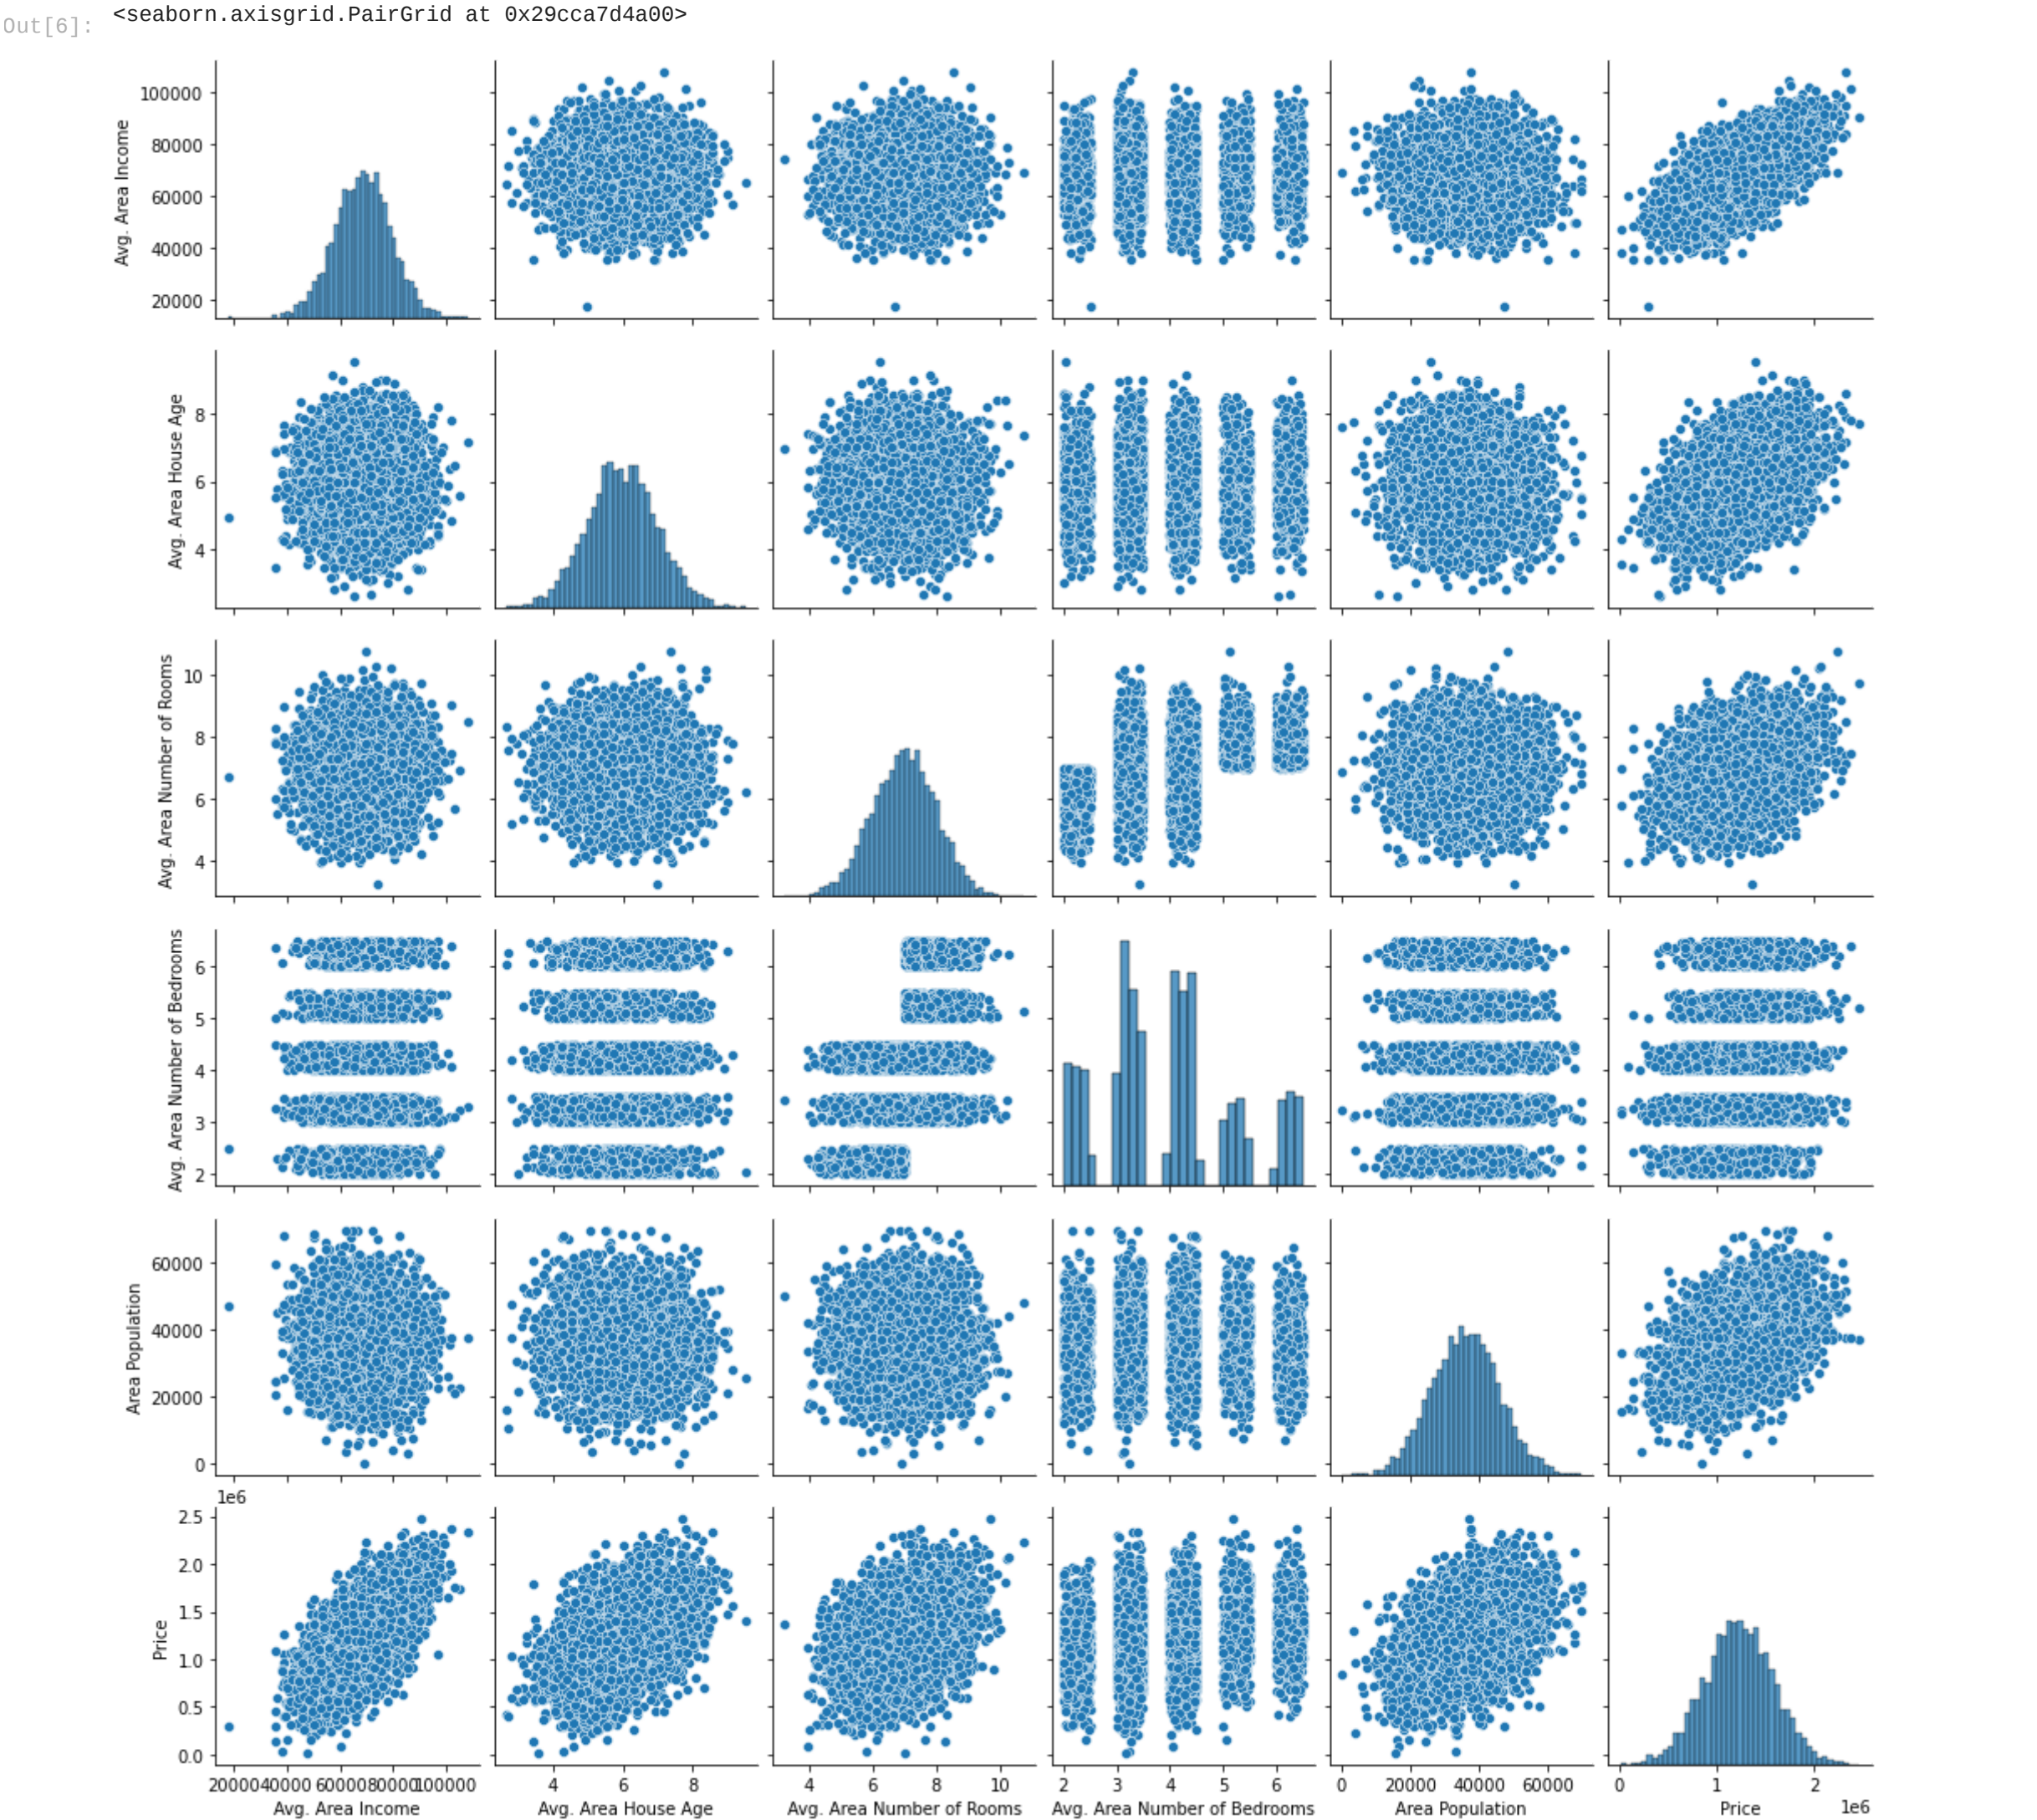
Out[4]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.45857	5.682861	7.009188	4.09	23086.80050	1059033.558	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.64245	6.002900	6.730821	3.09	40173.07217	1505890.915	188 Johnson Views Suite 079\nLake Kathleen, CA...

```
In [5]: #Determining the number of entries
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                      5000 non-null   float64
1   Avg. Area House Age                   5000 non-null   float64
2   Avg. Area Number of Rooms             5000 non-null   float64
3   Avg. Area Number of Bedrooms          5000 non-null   float64
4   Area Population                       5000 non-null   float64
5   Price                                 5000 non-null   float64
6   Address                              5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

```
In [6]: sns.pairplot(df)
```



```
In [7]: df.columns
```

Out[7]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms', 'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'], dtype='object')

```
In [8]: #Declaring the data into required feature and prediction feature
X = df[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms', 'Avg. Area Number of Bedrooms', 'Area Population']]
y = df['Price']
```

```
In [9]: #Splitting the data into training and testing set
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
In [10]: #Importing LinearRegression model
from sklearn.linear_model import LinearRegression
```

```
In [11]: #Initializing the model for use
lr = LinearRegression()
```

```
In [12]: #Fitting the training data on our model
lr.fit(X_train, y_train)
```

Out[12]: LinearRegression()

```
In [13]: #Intercept ----> Detrmining the value of response if the predictor variables are zero
print(lr.intercept_)

-2630108.9848230053
```

```
In [17]: #Coefficients ----> This determines that if all other features are held constant then 1 unit change in that feature will affect the price ogf target
lr.coef_
cdf = pd.DataFrame(lr.coef_, X.columns, columns = ['Coeff'])
cdf
```

Out[17]:

	Coeff
Avg. Area Income	21.543042
Avg. Area House Age	164974.342435
Avg. Area Number of Rooms	119936.487959
Avg. Area Number of Bedrooms	2163.200874
Area Population	15.260811

```
In [19]: #Predicting the prices
predictions = lr.predict(X_test)
```

```
In [20]: #Plotting the graph
plt.scatter(y_test, predictions)
```



```
In [21]: #Evaluation metrics for our model
from sklearn import metrics
```

```
In [28]: #Since this is a regression model we will determine the performance of our model on the basis of mae , mse rmse
mae = metrics.mean_absolute_error(y_test, predictions)
mse = metrics.mean_squared_error(y_test, predictions)
rmse = np.sqrt(metrics.mean_squared_error(y_test, predictions))
```

In [27]: print(mae)

81051.1958528423

In [29]: print(rmse)

101463.76909377535

In []: